

Veer Sanyal

Sammanish, WA · 360-643-9424 · vsanyal@purdue.edu · linkedin.com/in/veersanyal

Education

Purdue University, Mitch Daniels School of Business

West Lafayette, IN

B.S. Integrated Business and Engineering (IBE), Expected May 2028

Aug. 2024 – Present

- Selective business + engineering hybrid degree combining CS/data coursework with product and statistics training. Coursework: Statistics, Engineering Design & Project Execution, Managerial Accounting.

Projects

STICK: AI-Powered Study Platform

Independent product build, sole developer (React, TypeScript, Supabase)

2026 – Present

- Designed and shipped a structured, multi-step LLM pipeline (extract → classify → grade, on Deno edge functions) that turns raw exam PDFs into verified, structured question data: deliberately bounded steps, not an open-ended chatbot.
- Ingested a real 14-page, 12-question exam end-to-end with 100% extraction confidence; all 12 auto-proposed answer keys matched the instructor's published key on first pass, and every AI answer is cross-checked against an independent source before reaching a student.
- Designed the Postgres data model and access-control layer (46 migrations, 33+ row-level-security policies) and built the full user-facing loop end-to-end (auth, an FSRs-scheduled daily plan, and a server-graded question player), verified via an automated test.
- Run the build itself as a multi-agent workflow: a coordinator agent makes decisions serially and dispatches parallel agent lanes (code, design, research) into isolated git worktrees, governed by a 400+-entry decision ledger.

Cairn: Open-Source Scaffold for Agentic AI Systems

github.com/veer-sanyal/cairn

Claude Code plugin, sole author

2026

- Built and open-sourced a plugin that designs long-lived agentic systems around product fundamentals: it interviews the user to define a North Star metric, metric tree, and guardrails, then scaffolds a personalized multi-agent system with tiered memory, local telemetry, and human-gated self-improvement.
- Encoded 25 evidence-graded principles on agentic-system failure modes (context rot, objective corruption, knowledge expiry) with a bundled research engine using adversarial claim verification; adopted unprompted by an external builder who forked and extended it.

Experience

Firmly

Remote

Product and Engineering Intern

Jun. 2025 – Aug. 2025

- Shipped a PayPal checkout failure-recovery flow end-to-end: defined requirements and partnered with senior engineers to launch a guided recovery experience, replacing a generic error page, across ~15 high-risk payment edge cases.
- Built and tracked product metrics (error rate, recovery rate) via server-side logging (Node.js/Express) to drive data-driven decisions; built a 20+ test-case QA workflow adopted as the team standard.

Leadership & Activities

Engineering Projects in Community Service (EPICS)

West Lafayette, IN

Communications Officer & Build Lead

Aug. 2025 – Dec. 2025

- Led a 7-person cross-functional team building a rotating museum billboard prototype; translated requirements into build specs and presented design tradeoffs in two formal reviews to faculty, sponsors, and museum stakeholders.

Skills

AI & Data: agentic-system design & multi-agent orchestration (Claude Code), LLM pipeline / prompt design, evaluation design (AI-output cross-checking), generative-AI integration (Gemini API), SQL / Postgres data modeling, row-level access control

Tools: Python, HTML/CSS, Git/GitHub; React / TypeScript, Node.js / Express (built with); Excel, PowerPoint